

Introduction to Sockets Programming in C using TCP/IP

Introduction

- Computer Network
 - hosts, routers, communication channels
- **Hosts** run applications
- **Routers** forward information
- **Packets**: sequence of bytes
 - contain control information
 - e.g. destination host
- **Protocol** is an agreement
 - meaning of packets
 - structure and size of packetse.g. Hypertext Transfer Protocol (HTTP)




Protocol Families - TCP/IP

- Several protocols for different problems
- ☞ **Protocol Suites** or **Protocol Families**: TCP/IP
- TCP/IP provides **end-to-end** connectivity specifying how data should be
 - formatted,
 - addressed,
 - transmitted,
 - routed, and
 - received at the destination
- can be used in the internet and in stand-alone private networks
- it is organized into **layers**

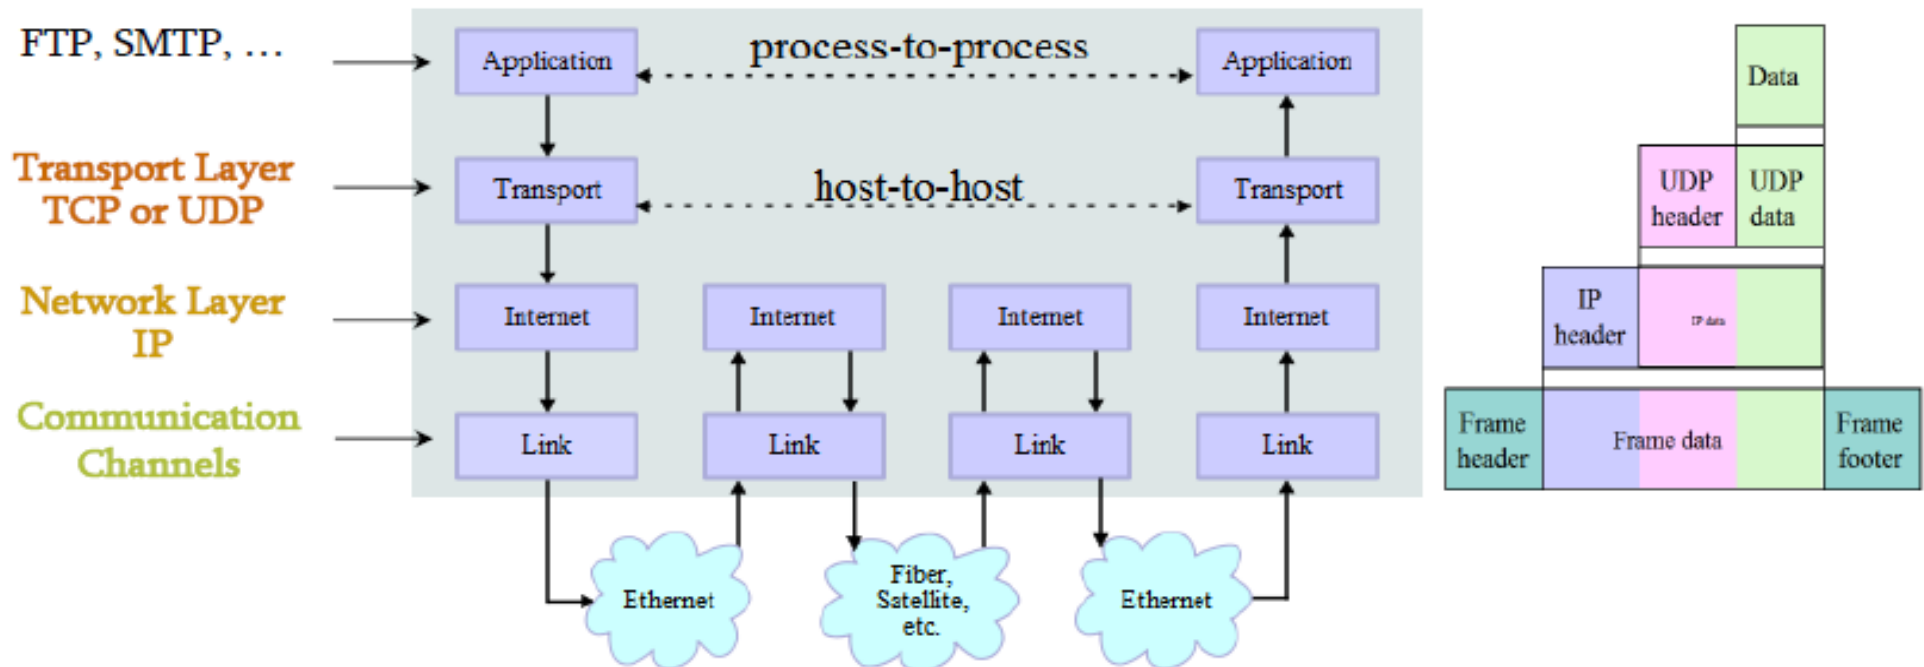
TCP/IP

Network Topology



*

Data Flow



* Image is taken from "http://en.wikipedia.org/wiki/TCP/IP_model"

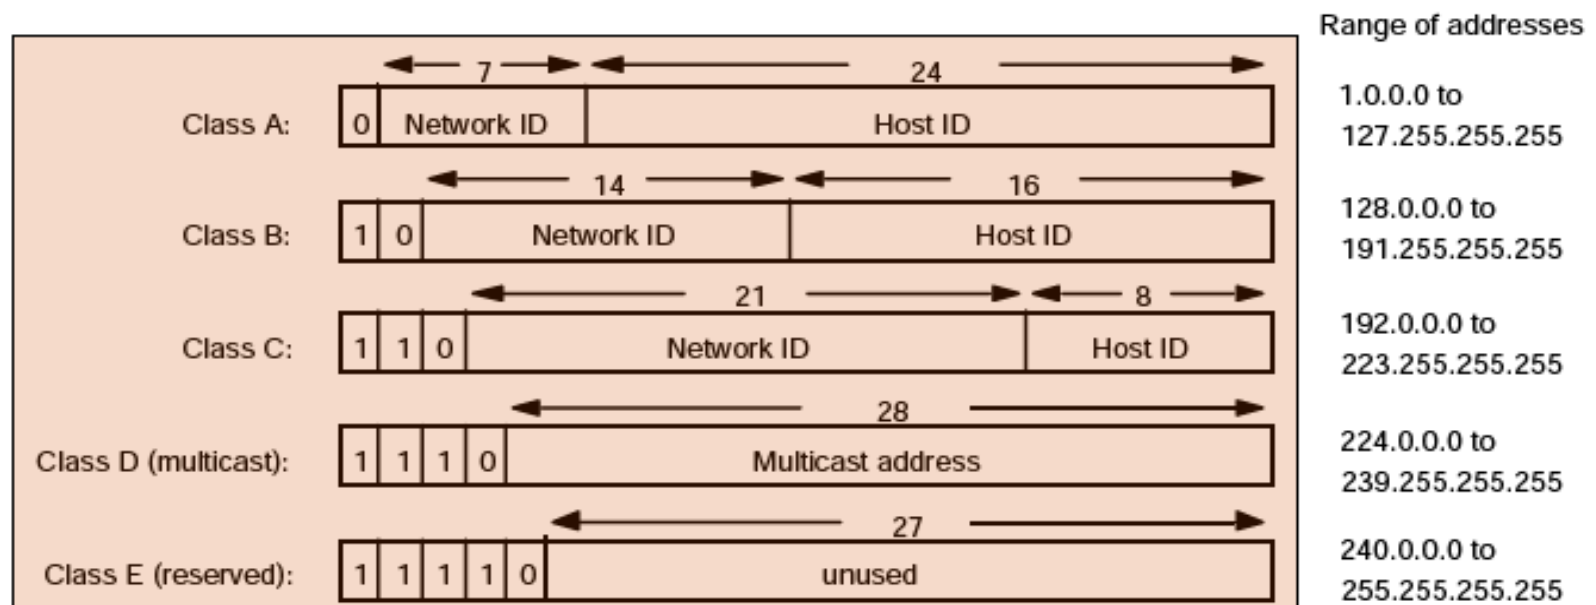
Internet Protocol (IP)

- provides a **datagram** service
 - packets are handled and delivered independently
- **best-effort** protocol
 - may lose, reorder or duplicate packets
- each packet must contain an **IP address** of its destination

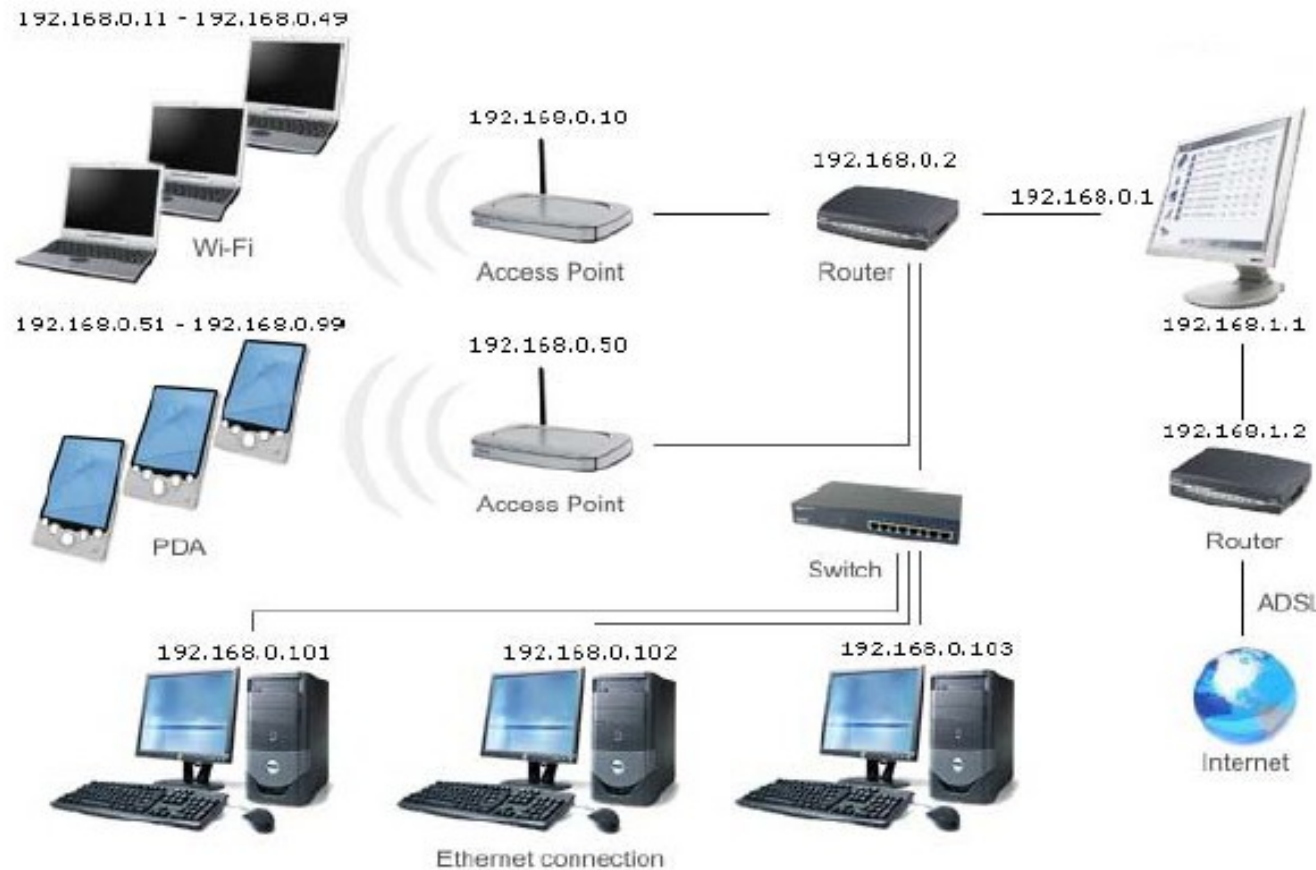


Addresses - IPv4

- The **32** bits of an IPv4 address are broken into **4 octets**, or 8 bit fields (0-255 value in decimal notation).
- For networks of different size,
 - the first one (for large networks) to three (for small networks) octets can be used to identify the **network**, while
 - the rest of the octets can be used to identify the **node** on the network.



Local Area Network Addresses - IPv4



TCP vs UDP

- Both use **port numbers**
 - application-specific construct serving as a communication endpoint
 - 16-bit unsigned integer, thus ranging from 0 to 65535
 - ☞ to provide **end-to-end** transport
- UDP: User Datagram Protocol
 - no acknowledgements
 - no retransmissions
 - out of order, duplicates possible
 - connectionless, i.e., app indicates destination for each packet
- TCP: Transmission Control Protocol
 - reliable **byte-stream channel** (in order, all arrive, no duplicates)
 - similar to file I/O
 - flow control
 - connection-oriented
 - bidirectional


TCP vs UDP

- TCP is used for services with a large data capacity, and a persistent connection
- UDP is more commonly used for quick lookups, and single use query-reply actions.
- Some common examples of TCP and UDP with their default ports:

DNS lookup	UDP	53
FTP	TCP	21
HTTP	TCP	80
POP3	TCP	110
Telnet	TCP	23

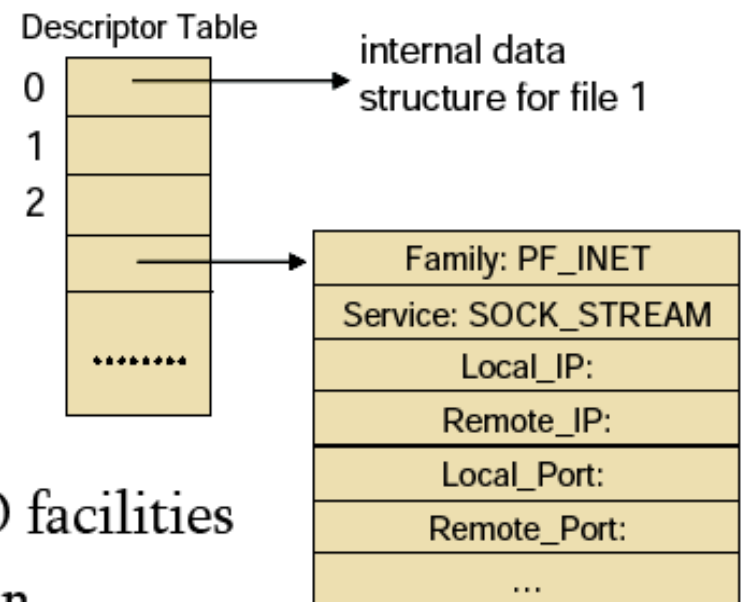
Berkley Sockets

- Universally known as **Sockets**
- It is an abstraction through which an application may send and receive data
- Provide **generic access** to interprocess communication services
 - e.g. IPX/SPX, Appletalk, TCP/IP
- Standard API for networking

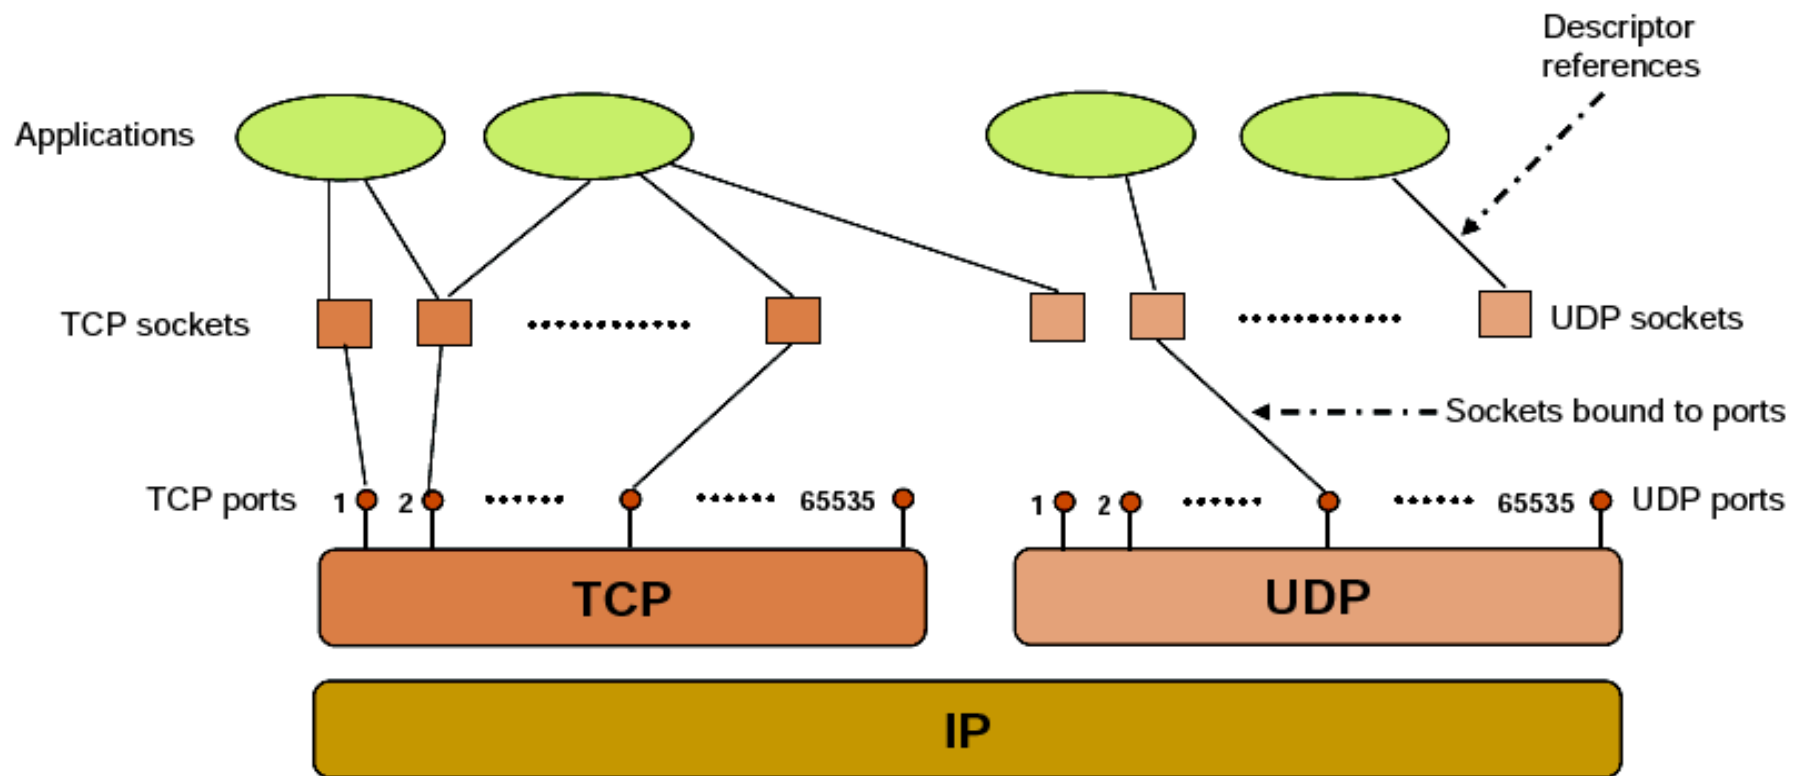


Sockets

- Uniquely identified by
 - an internet address
 - an end-to-end protocol (e.g. TCP or UDP)
 - a port number
- Two types of (TCP/IP) sockets
 - **Stream** sockets (e.g. uses TCP)
 - provide reliable byte-stream service
 - **Datagram** sockets (e.g. uses UDP)
 - provide best-effort datagram service
 - messages up to 65.500 bytes
- Socket extend the convectional UNIX I/O facilities
 - file descriptors for network communication
 - extended the read and write system calls



Sockets



Socket Programming



Client-Server communication

■ Server

- passively waits for and responds to clients
- **passive** socket

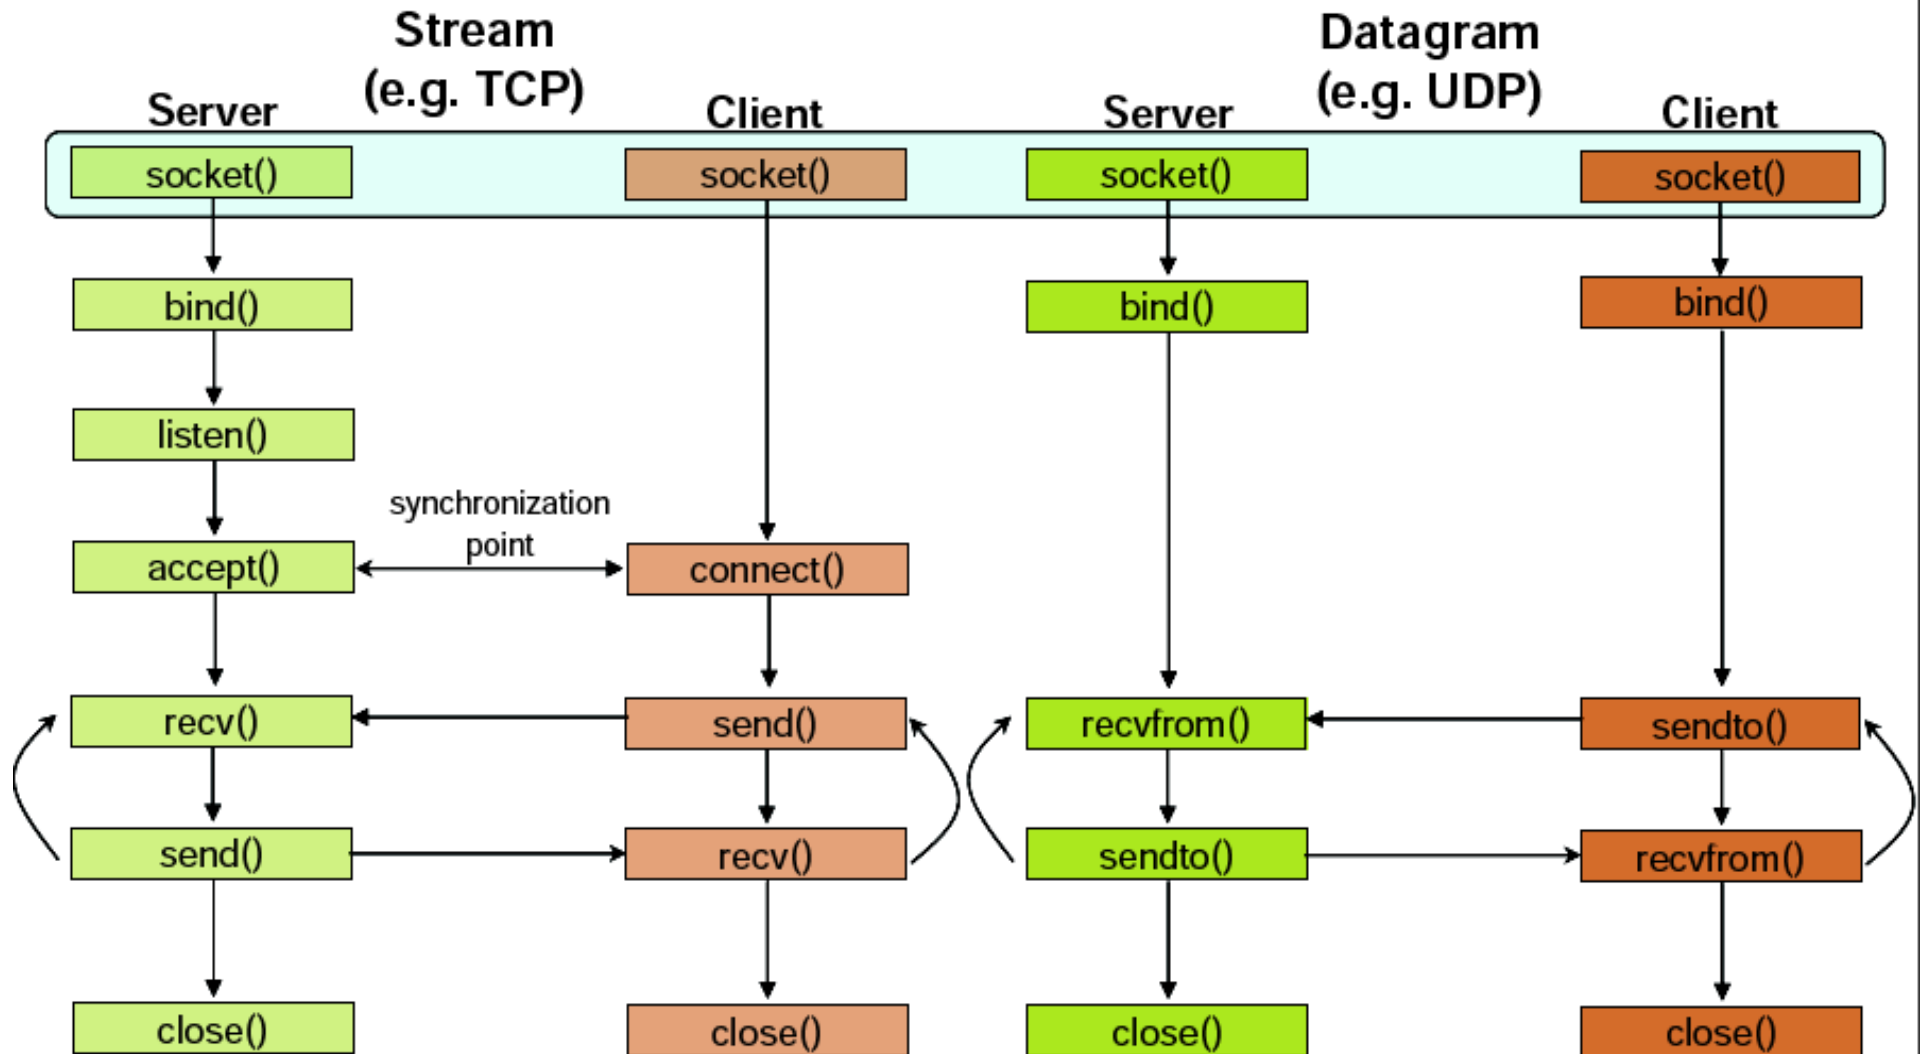
■ Client

- initiates the communication
- must know the address and the port of the server
- **active** socket

Sockets - Procedures

Primitive	Meaning
Socket	Create a new communication endpoint
Bind	Attach a local address to a socket
Listen	Announce willingness to accept connections
Accept	Block caller until a connection request arrives
Connect	Actively attempt to establish a connection
Send	Send some data over the connection
Receive	Receive some data over the connection
Close	Release the connection

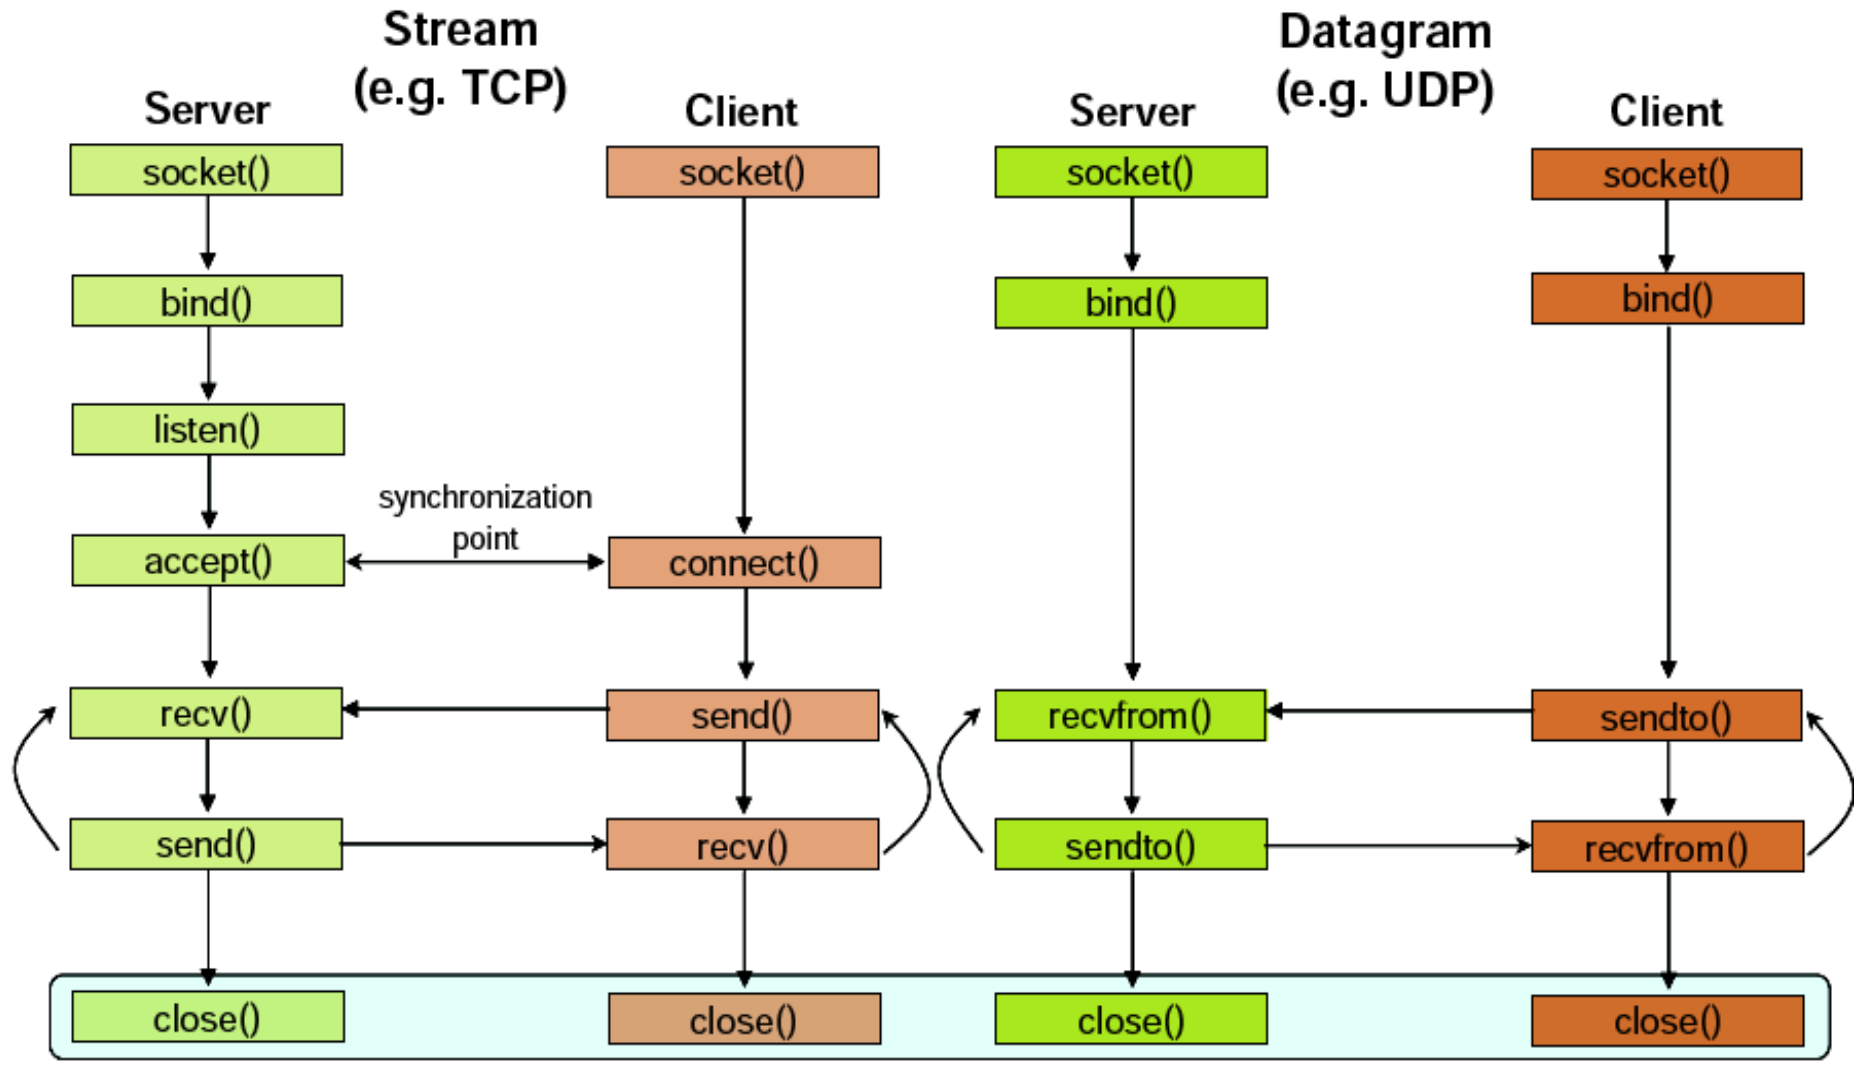
Client - Server Communication - Unix



Socket creation in C: `socket()`

- `int sockid = socket(family, type, protocol);`
 - **sockid**: socket descriptor, an integer (like a file-handle)
 - **family**: integer, communication domain, e.g.,
 - `PF_INET`, IPv4 protocols, Internet addresses (typically used)
 - `PF_UNIX`, Local communication, File addresses
 - **type**: communication type
 - `SOCK_STREAM` - reliable, 2-way, connection-based service
 - `SOCK_DGRAM` - unreliable, connectionless, messages of maximum length
 - **protocol**: specifies protocol
 - `IPPROTO_TCP` `IPPROTO_UDP`
 - usually set to 0 (i.e., use default protocol)
 - upon failure returns -1
- ☞ NOTE: socket call does not specify where data will be coming from, nor where it will be going to – it just creates the interface!

Client - Server Communication - Unix



Socket close in C: `close()`

- When finished using a socket, the socket should be closed
- `status = close(sockid);`
 - **sockid**: the file descriptor (socket being closed)
 - **status**: 0 if successful, -1 if error
- Closing a socket
 - closes a connection (for stream socket)
 - frees up the port used by the socket

Specifying Addresses

- Socket API defines a **generic** data type for addresses:

```
struct sockaddr {  
    unsigned short sa_family; /* Address family (e.g. AF_INET) */  
    char sa_data[14];        /* Family-specific address information */  
}
```

- Particular form of the sockaddr used for **TCP/IP** addresses:

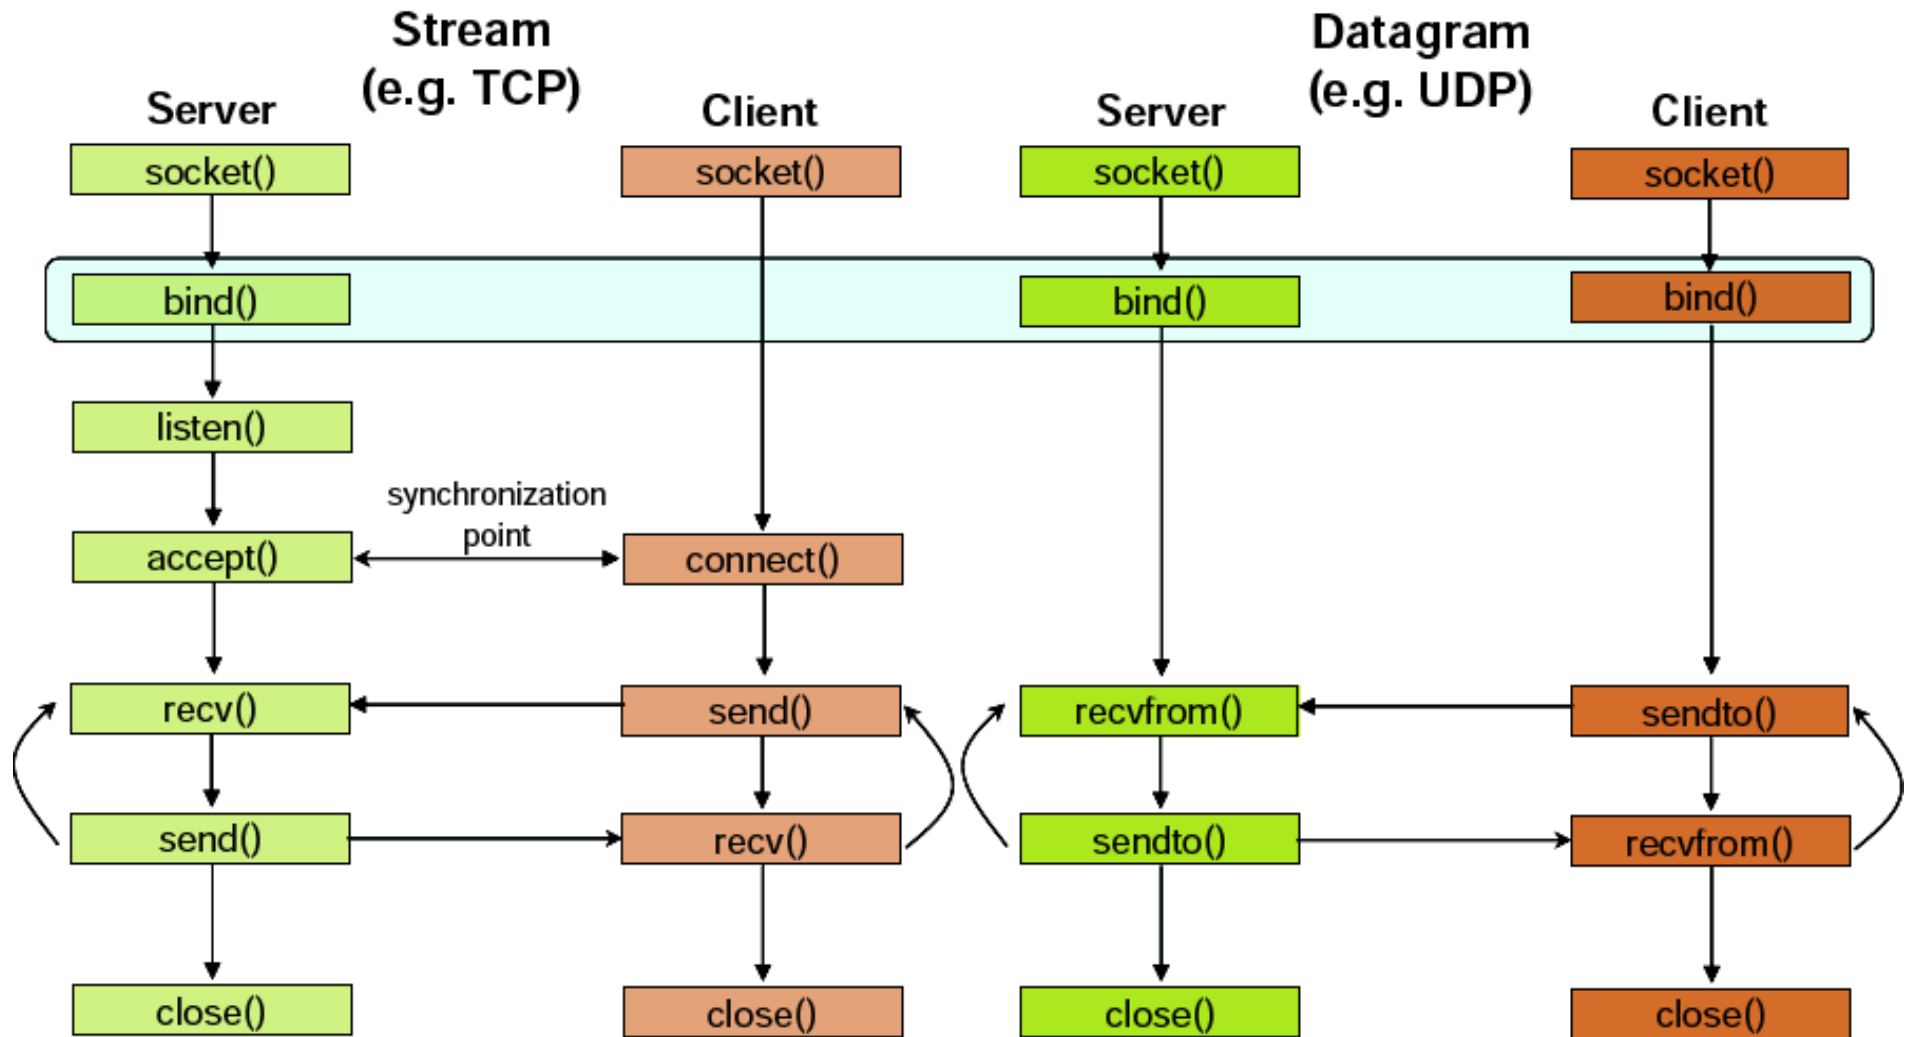
```
struct in_addr {  
    unsigned long s_addr; /* Internet address (32 bits) */  
}  
  
struct sockaddr_in {  
    unsigned short sin_family; /* Internet protocol (AF_INET) */  
    unsigned short sin_port;   /* Address port (16 bits) */  
    struct in_addr sin_addr;   /* Internet address (32 bits) */  
    char sin_zero[8];         /* Not used */  
}
```

👉 **Important:** sockaddr_in can be casted to a sockaddr

Specifying Addresses

	sa_family		sa_data	
sockaddr	Family	Blob (14 bytes)		
	2 bytes	2 bytes	4 bytes	8 bytes
sockaddr_in	Family	Port	Internet address	Unused
	sin_family	sin_port	sin_addr	sin_zero

Client - Server Communication - Unix



Assign address to socket: `bind()`

- associates and reserves a port for use by the socket
- `int status = bind(sockid, &addrport, size);`
 - **sockid**: integer, socket descriptor
 - **addrport**: struct `sockaddr`, the (IP) address and port of the machine
 - for TCP/IP server, internet address is usually set to `INADDR_ANY`, i.e., chooses any incoming interface
 - **size**: the size (in bytes) of the `addrport` structure
 - **status**: upon failure -1 is returned

bind() - Example with TCP

```
int sockid;
struct sockaddr_in addrport;
sockid = socket(PF_INET, SOCK_STREAM, 0);

addrport.sin_family = AF_INET;
addrport.sin_port = htons(5100);
addrport.sin_addr.s_addr = htonl(INADDR_ANY);
if(bind(sockid, (struct sockaddr *) &addrport, sizeof(addrport)) != -1) {
    ...}
```

Skipping the bind ()

- bind can be skipped for both types of sockets
- Datagram socket:
 - if only sending, no need to bind. The OS finds a port each time the socket sends a packet
 - if receiving, need to bind
- Stream socket:
 - destination determined during connection setup
 - don't need to know port sending from (during connection setup, receiving end is informed of port)

Exchanging data with stream socket

- `int count = send(sockid, msg, msgLen, flags);`
 - `msg`: `const void[]`, message to be transmitted
 - `msgLen`: integer, length of message (in bytes) to transmit
 - `flags`: integer, special options, usually just 0
 - `count`: # bytes transmitted (-1 if error)
- `int count = recv(sockid, recvBuf, bufLen, flags);`
 - `recvBuf`: `void[]`, stores received bytes
 - `bufLen`: # bytes received
 - `flags`: integer, special options, usually just 0
 - `count`: # bytes received (-1 if error)
- Calls are **blocking**
 - returns only after data is sent / received

Exchanging data with datagram socket

- `int count = sendto(sockid, msg, msgLen, flags, &foreignAddr, addrlen);`
 - `msg, msgLen, flags, count`: same with `send()`
 - `foreignAddr`: struct `sockaddr`, address of the destination
 - `addrlen`: `sizeof(foreignAddr)`
- `int count = recvfrom(sockid, recvBuf, bufLen, flags, &clientAddr, addrlen);`
 - `recvBuf, bufLen, flags, count`: same with `recv()`
 - `clientAddr`: struct `sockaddr`, address of the client
 - `addrlen`: `sizeof(clientAddr)`
- Calls are **blocking**
 - returns only after data is sent / received


Constructing Messages - Byte Ordering

- Address and port are stored as integers

- u_short sin_port; (16 bit)
- in_addr sin_addr; (32 bit)

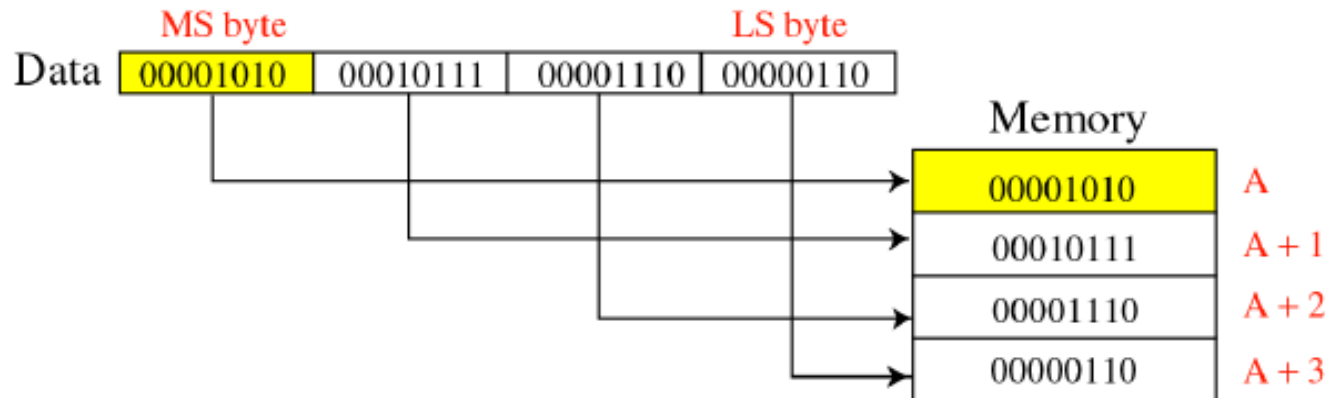
□ Problem:

- different machines / OS's use different word orderings
 - little-endian: lower bytes first
 - big-endian: higher bytes first
- these machines may communicate with one another over the network

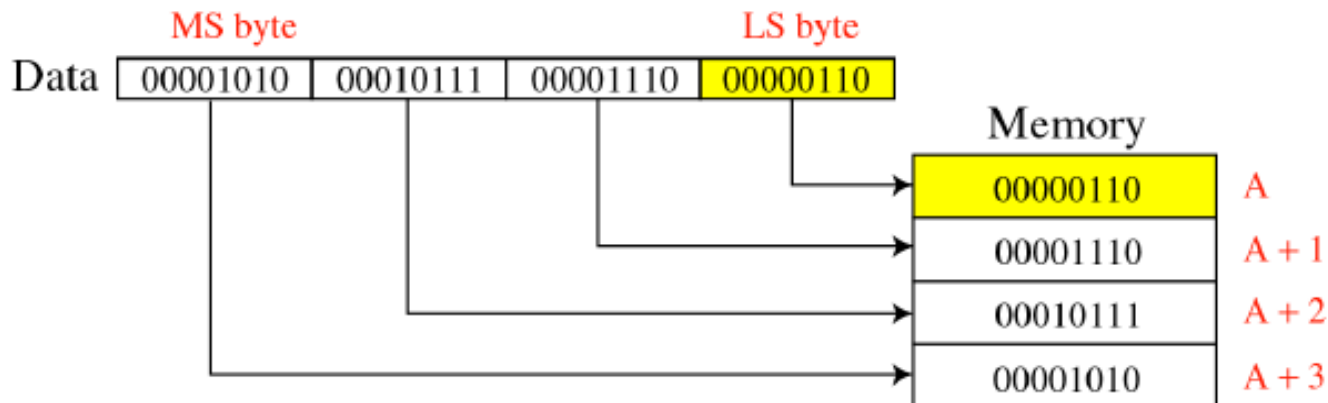


Constructing Messages - Byte Ordering

- Big-Endian:




- Little-Endian:



Constructing Messages - Byte Ordering - Solution: Network Byte Ordering

- **Host Byte-Ordering**: the byte ordering used by a host (big or little)
 - **Network Byte-Ordering**: the byte ordering used by the network – always big-endian
 - `u_long htonl(u_long x);` ■ `u_long ntohl(u_long x);`
 - `u_short htons(u_short x);` ■ `u_short ntohs(u_short x);`
- On big-endian machines, these routines do nothing
 - On little-endian machines, they reverse the byte order



Constructing Messages - Byte Ordering - Example

Client

```
unsigned short clientPort, message;    unsigned int messageLength;

servPort = 1111;
message = htons(clientPort);
messageLength = sizeof(message);

if (sendto( clientSock, message, messageLength, 0,
            (struct sockaddr *) &echoServAddr, sizeof(echoServAddr))
    != messageLength)
    DieWithError("send() sent a different number of bytes than expected");
```

Server

```
unsigned short clientPort, rcvBuffer;
unsigned int rcvMsgSize ;

if ( recvfrom(servSock, &rcvBuffer, sizeof(unsigned int), 0),
    (struct sockaddr *) &echoClientAddr, sizeof(echoClientAddr)) < 0)
    DieWithError("recvfrom() failed");

clientPort = ntohs(rcvBuffer);
printf ("Client's port: %d", clientPort);
```

Example - Echo

- A client communicates with an “echo” server
- The server simply echoes whatever it receives back to the client

Example - Echo using datagram socket

```
/* Create socket for sending/receiving datagrams */  
if ((servSock = socket(PF_INET, SOCK_DGRAM, IPPROTO_UDP)) < 0)  
    DieWithError("socket() failed");
```

```
/* Create a datagram/UDP socket */  
if ((clientSock = socket(PF_INET, SOCK_DGRAM, IPPROTO_UDP)) < 0)  
    DieWithError("socket() failed");
```

Client

1. **Create a UDP socket**
2. Assign a port to socket
3. Communicate
4. Close the socket

Server

1. **Create a UDP socket**
2. Assign a port to socket
3. Repeatedly
 - Communicate

Example - Echo using datagram socket

```
echoServAddr.sin_family = AF_INET;           /* Internet address family */
echoServAddr.sin_addr.s_addr = htonl(INADDR_ANY); /* Any incoming interface */
echoServAddr.sin_port = htons(echoServPort);  /* Local port */

if (bind(servSock, (struct sockaddr *) &echoServAddr, sizeof(echoServAddr)) < 0)
    DieWithError("bind() failed");
```

```
echoClientAddr.sin_family = AF_INET;           /* Internet address family */
echoClientAddr.sin_addr.s_addr = htonl(INADDR_ANY); /* Any incoming interface */
echoClientAddr.sin_port = htons(echoClientPort); /* Local port */

if(bind(clientSock, (struct sockaddr *) &echoClientAddr, sizeof(echoClientAddr)) < 0)
    DieWithError("connect() failed");
```

Client

1. Create a UDP socket
2. **Assign a port to socket**
3. Communicate
4. Close the socket

Server

1. Create a UDP socket
2. **Assign a port to socket**
3. Repeatedly
 - Communicate

Example - Echo using datagram socket

```
echoServAddr.sin_family = AF_INET;           /* Internet address family */
echoServAddr.sin_addr.s_addr = inet_addr(echoservIP); /* Server IP address*/
echoServAddr.sin_port = htons(echoServPort); /* Server port */

echoStringLength = strlen(echoString); /* Determine input length */

/* Send the string to the server */
if (sendto( clientSock, echoString, echoStringLength, 0,
           (struct sockaddr *) &echoServAddr, sizeof(echoServAddr))
    != echoStringLength)
    DieWithError("send() sent a different number of bytes than expected");
```

Client

1. Create a UDP socket
2. Assign a port to socket
3. **Communicate**
4. Close the socket

Server

1. Create a UDP socket
2. **Assign a port to socket**
3. Repeatedly
 - Communicate

Example - Echo using datagram socket

```
for (;;) /* Run forever */
{
    clientAddrLen = sizeof(echoClientAddr) /* Set the size of the in-out parameter */
    /*Block until receive message from client*/
    if ((recvMsgSize = recvfrom(servSock, echoBuffer, ECHOMAX, 0),
        (struct sockaddr *) &echoClientAddr, sizeof(echoClientAddr))) < 0)
        DieWithError("recvfrom() failed");

    if (sendto(servSock, echobuffer, recvMsgSize, 0,
        (struct sockaddr *) &echoClientAddr, sizeof(echoClientAddr))
        != recvMsgSize)
        DieWithError("send() failed");
}
```

Client

1. Create a UDP socket
2. Assign a port to socket
3. **Communicate**
4. Close the socket

Server

1. Create a UDP socket
2. Assign a port to socket
3. Repeatedly
 - **Communicate**

Example - Echo using datagram socket

Similarly, the client receives the data from the server

Client

1. Create a UDP socket
2. Assign a port to socket
3. **Communicate**
4. Close the socket

Server

1. Create a UDP socket
2. Assign a port to socket
3. Repeatedly
 - **Communicate**

Example - Echo using datagram socket

```
close(clientSock);
```


Client

1. Create a UDP socket
2. Assign a port to socket
3. Communicate
4. **Close the socket**

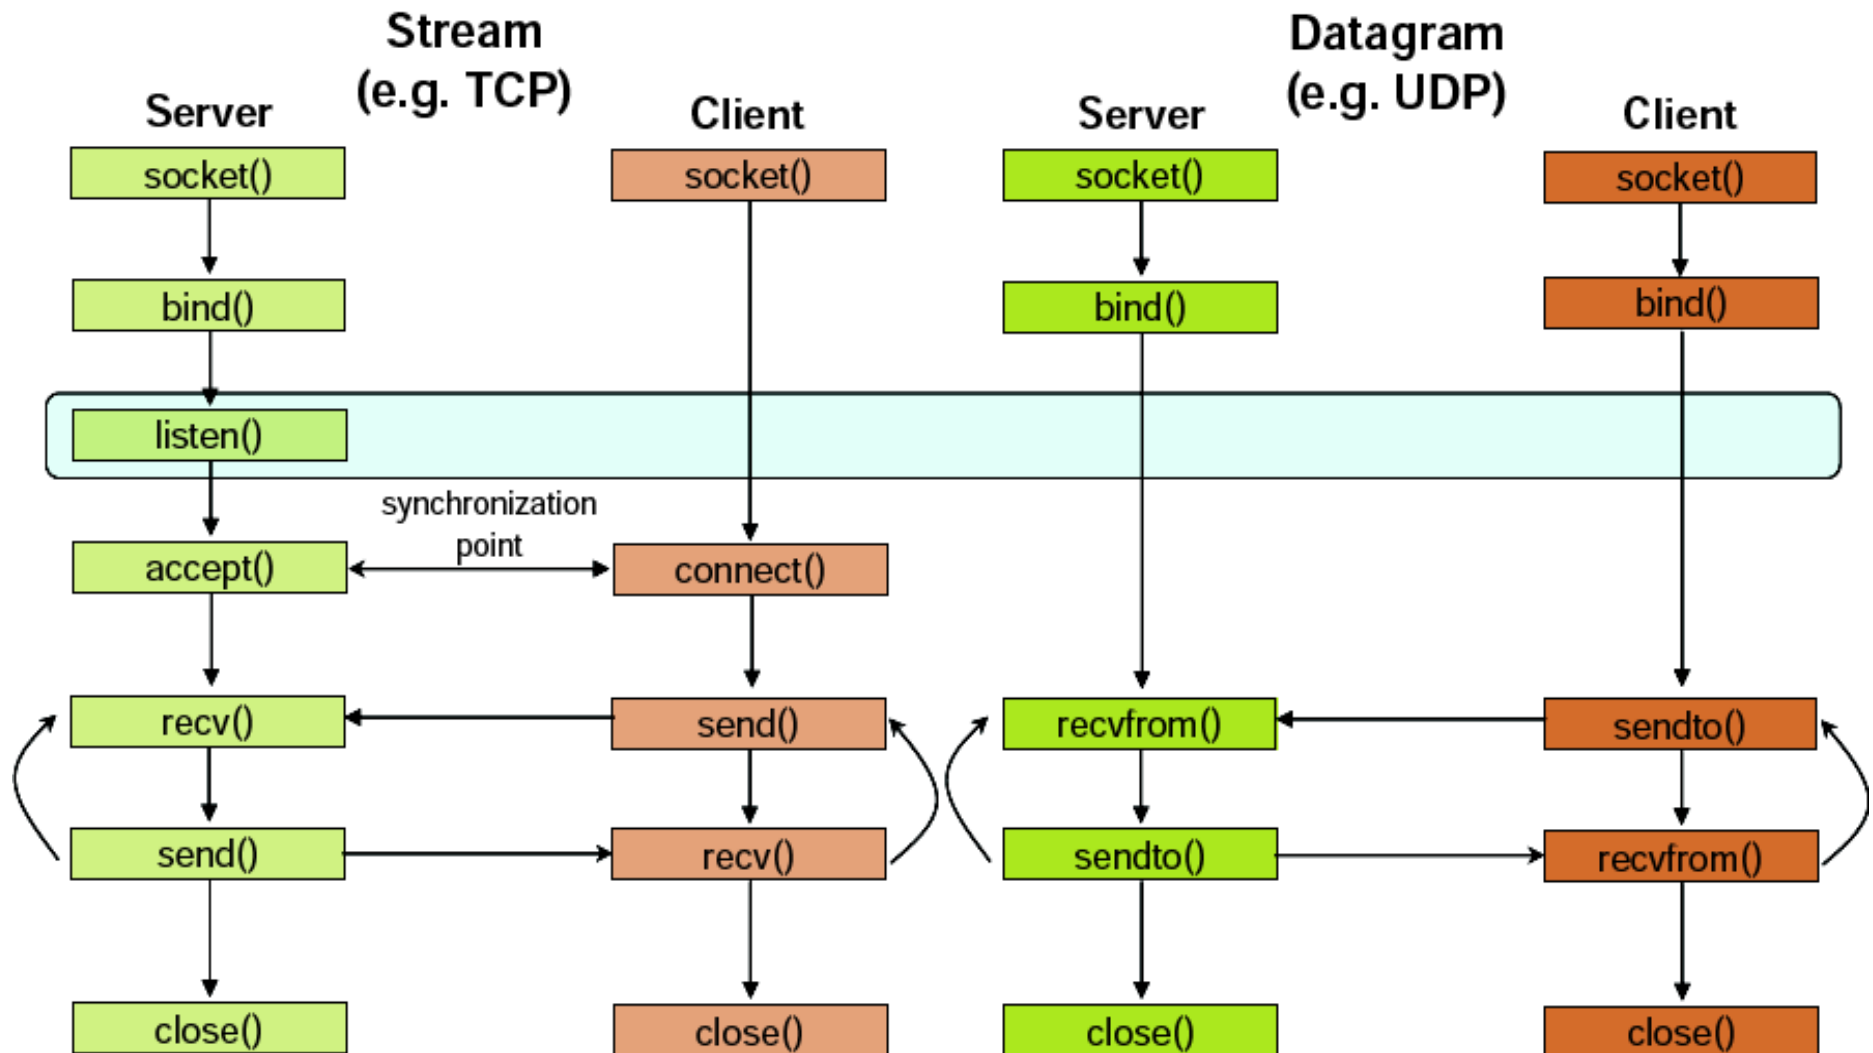
Server

1. Create a UDP socket
2. Assign a port to socket
3. **Repeatedly**
 - Communicate

Client - Server Communication - Unix



Client - Server Communication - Unix

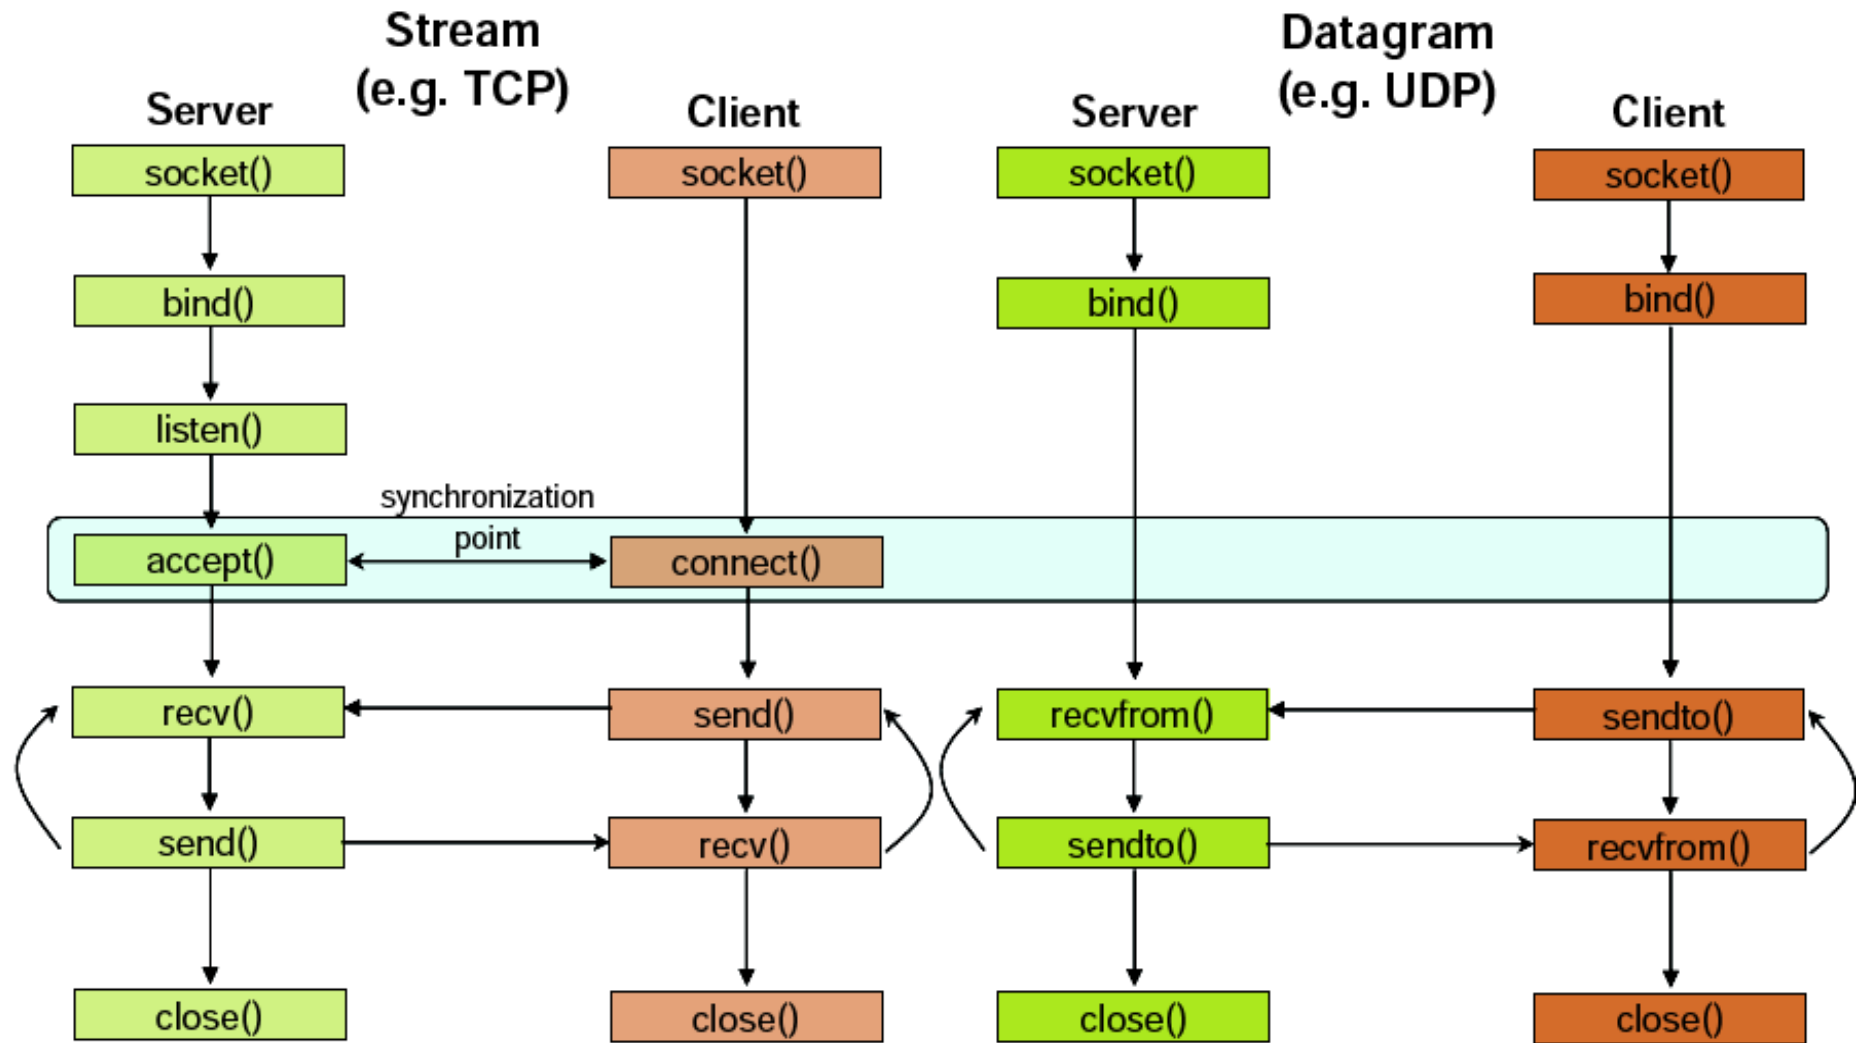


Assign address to socket: `bind()`

- Instructs TCP protocol implementation to listen for connections
- ```
int status = listen(sockid, queueLimit);
```

  - **sockid**: integer, socket descriptor
  - **queueLen**: integer, # of active participants that can “wait” for a connection
  - **status**: 0 if listening, -1 if error
- `listen()` is **non-blocking**: returns immediately
- The listening socket (`sockid`)
  - is never used for sending and receiving
  - is used by the server only as a way to get new sockets

# Client - Server Communication - Unix



# Establish Connection: connect ( )

- The client establishes a connection with the server by calling connect ( )
- ```
int status = connect(sockid, &foreignAddr, addrlen);
```

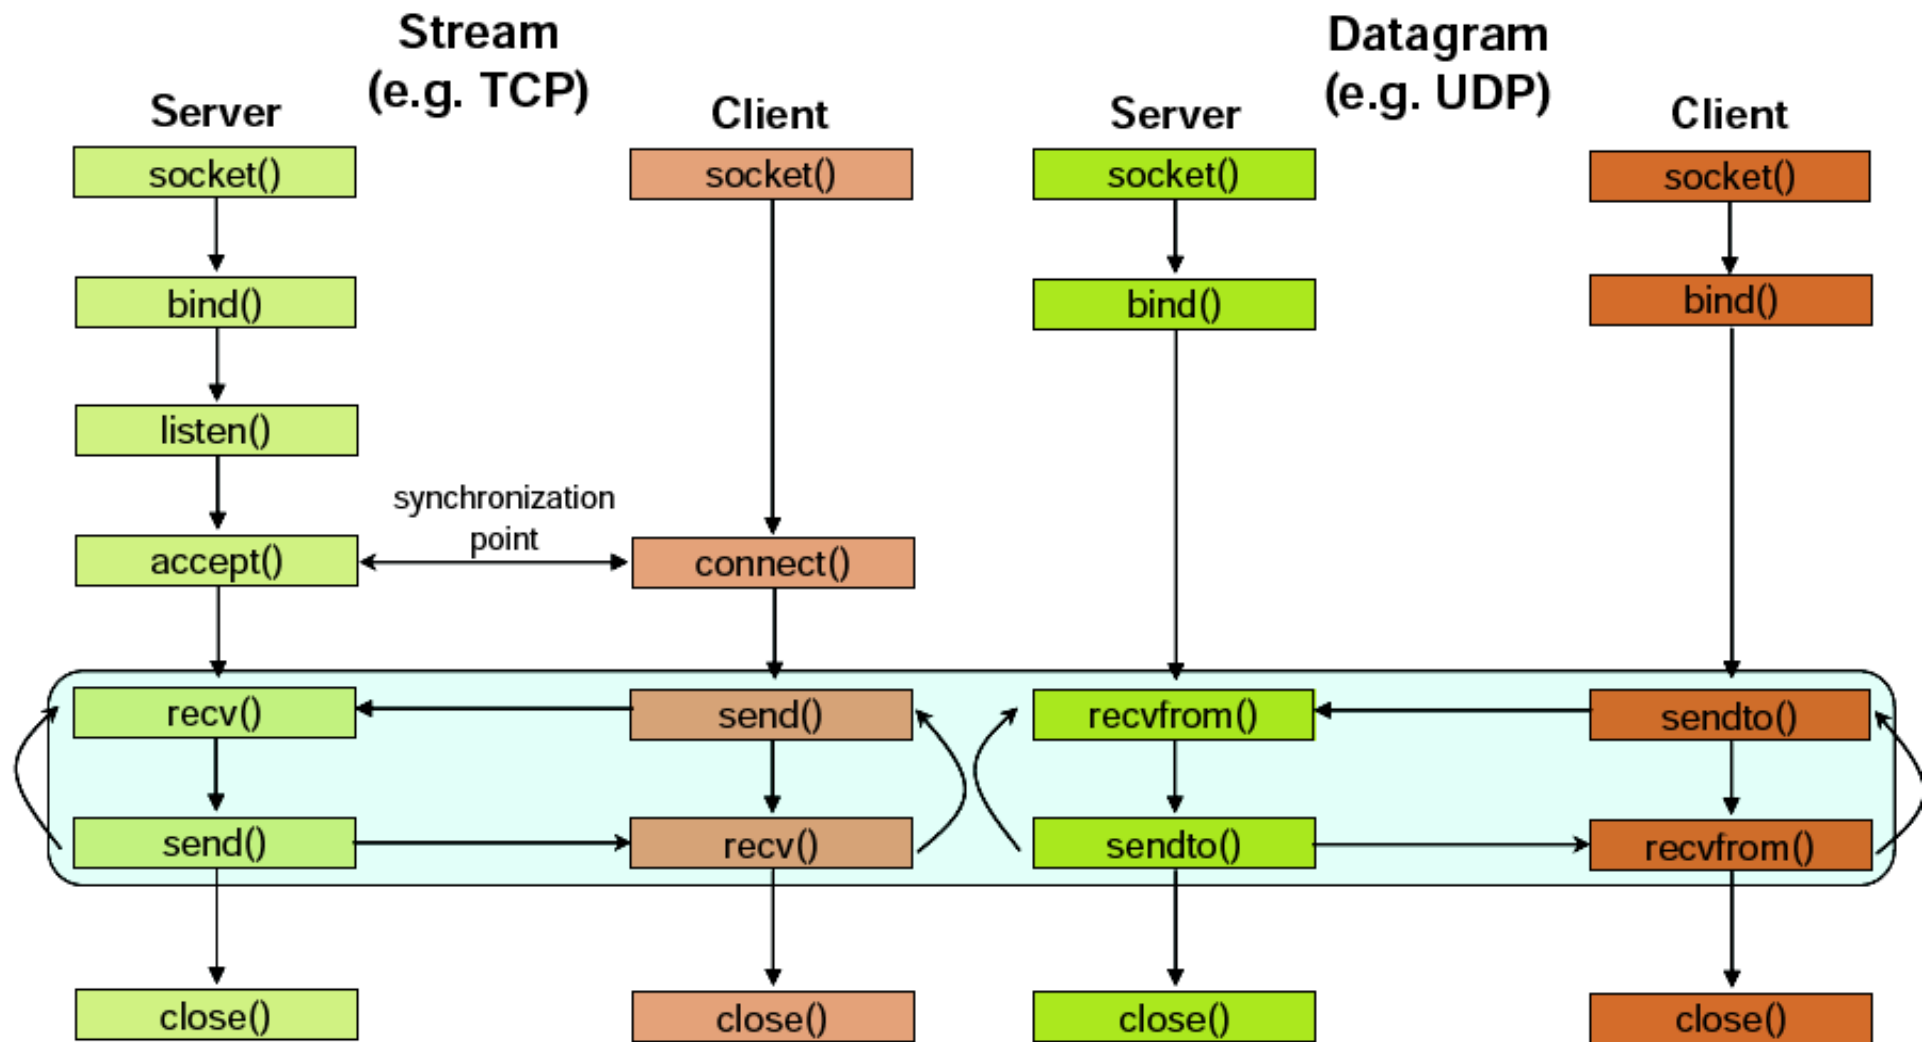
 - **sockid**: integer, socket to be used in connection
 - **foreignAddr**: struct sockaddr: address of the passive participant
 - **addrlen**: integer, sizeof(name)
 - status: 0 if successful connect, -1 otherwise
- connect () is **blocking**

Incoming Connection: `accept()`

- The server gets a socket for an incoming client connection by calling `accept()`
- ```
int s = accept(sockid, &clientAddr, &addrLen);
```

  - **s**: integer, the new socket (used for data-transfer)
  - **sockid**: integer, the orig. socket (being listened on)
  - **clientAddr**: struct `sockaddr`, address of the active participant
    - filled in upon return
  - **addrLen**: `sizeof(clientAddr)`: value/result parameter
    - must be set appropriately before call
    - adjusted upon return
- `accept()`
  - is **blocking**: waits for connection before returning
  - dequeues the next connection on the queue for socket (`sockid`)

# Client - Server Communication - Unix



# Exchanging data with stream socket

- `int count = send(sockid, msg, msgLen, flags);`
  - `msg`: `const void[]`, message to be transmitted
  - `msgLen`: integer, length of message (in bytes) to transmit
  - `flags`: integer, special options, usually just 0
  - `count`: # bytes transmitted (-1 if error)
- `int count = recv(sockid, recvBuf, bufLen, flags);`
  - `recvBuf`: `void[]`, stores received bytes
  - `bufLen`: # bytes received
  - `flags`: integer, special options, usually just 0
  - `count`: # bytes received (-1 if error)
- Calls are **blocking**
  - returns only after data is sent / received

---

# Example - Echo using stream socket

The server starts by getting ready to receive client connections...

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. Communicate
  - c. Close the connection

# Example - Echo using stream socket

```
/* Create socket for incoming connections */
if ((servSock = socket(PF_INET, SOCK_STREAM, IPPROTO_TCP)) < 0)
 DieWithError("socket() failed");
```

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. **Create a TCP socket**
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. Communicate
  - c. Close the connection



# Example - Echo using stream socket

```
echoServAddr.sin_family = AF_INET; /* Internet address family */
echoServAddr.sin_addr.s_addr = htonl(INADDR_ANY); /* Any incoming interface */
echoServAddr.sin_port = htons(echoServPort); /* Local port */

if (bind(servSock, (struct sockaddr *) &echoServAddr, sizeof(echoServAddr)) < 0)
 DieWithError("bind() failed");
```

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. **Assign a port to socket**
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. Communicate
  - c. Close the connection

# Example - Echo using stream socket

```
/* Mark the socket so it will listen for incoming connections */
if (listen(servSock, MAXPENDING) < 0)
 DieWithError("listen() failed");
```

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. **Set socket to listen**
4. Repeatedly:
  - a. Accept new connection
  - b. Communicate
  - c. Close the connection

# Example - Echo using stream socket

```
for (;;) /* Run forever */
{
 clntLen = sizeof(echoClntAddr);

 if ((clientSock=accept(servSock, (struct sockaddr *)&echoClntAddr, &clntLen)) < 0)
 DieWithError("accept() failed");
 ...
}
```

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection

# Example - Echo using stream socket

Server is now blocked waiting for connection from a client

...

A client decides to talk to the server

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection

# Example - Echo using stream socket

```
/* Create a reliable, stream socket using TCP */
if ((clientSock = socket(PF_INET, SOCK_STREAM, IPPROTO_TCP)) < 0)
 DieWithError("socket() failed");
```

## Client

1. **Create a TCP socket**
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection

# Example - Echo using stream socket

```
echoServAddr.sin_family = AF_INET; /* Internet address family */
echoServAddr.sin_addr.s_addr = inet_addr(echoServIP); /* Server IP address*/
echoServAddr.sin_port = htons(echoServPort); /* Server port */

if (connect(clientSock, (struct sockaddr *) &echoServAddr,
 sizeof(echoServAddr)) < 0)
 DieWithError("connect() failed");
```

## Client

1. Create a TCP socket
2. **Establish connection**
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection

# Example - Echo using stream socket

Server's accept procedure is now unblocked and returns client's socket

```
for (;;) /* Run forever */
{
 clntLen = sizeof(echoClntAddr);

 if ((clientSock=accept(servSock, (struct sockaddr *)&echoClntAddr,&clntLen))<0)
 DieWithError("accept() failed");
 ...
}
```

## Client

1. Create a TCP socket
2. **Establish connection**
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection



# Example - Echo using stream socket

```
echoStringLen = strlen(echoString); /* Determine input length */

/* Send the string to the server */
if (send(clientSock, echoString, echoStringLen, 0) != echoStringLen)
 DieWithError("send() sent a different number of bytes than expected");
```

## Client

1. Create a TCP socket
2. Establish connection
3. **Communicate**
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection



# Example - Echo using stream socket

```
/* Receive message from client */
if ((recvMsgSize = recv(clntSocket, echoBuffer, RCVBUFSIZE, 0)) < 0)
 DieWithError("recv() failed");
/* Send received string and receive again until end of transmission */
while (recvMsgSize > 0) { /* zero indicates end of transmission */
 if (send(clientSocket, echobuffer, recvMsgSize, 0) != recvMsgSize)
 DieWithError("send() failed");
 if ((recvMsgSize = recv(clientSocket, echoBuffer, RCVBUFSIZE, 0)) < 0)
 DieWithError("recv() failed");
}
```

## Client

1. Create a TCP socket
2. Establish connection
3. **Communicate**
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. **Communicate**
  - c. Close the connection

---

# Example - Echo using stream socket

Similarly, the client receives the data from the server

## Client

1. Create a TCP socket
2. Establish connection
3. **Communicate**
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. **Communicate**
  - c. Close the connection

# Example - Echo using stream socket

```
close(clientSock);
```

## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. **Close the connection**

```
close(clientSock);
```

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. Accept new connection
  - b. Communicate
  - c. **Close the connection**

---

# Example - Echo using stream socket

Server is now blocked waiting for connection from a client  
...

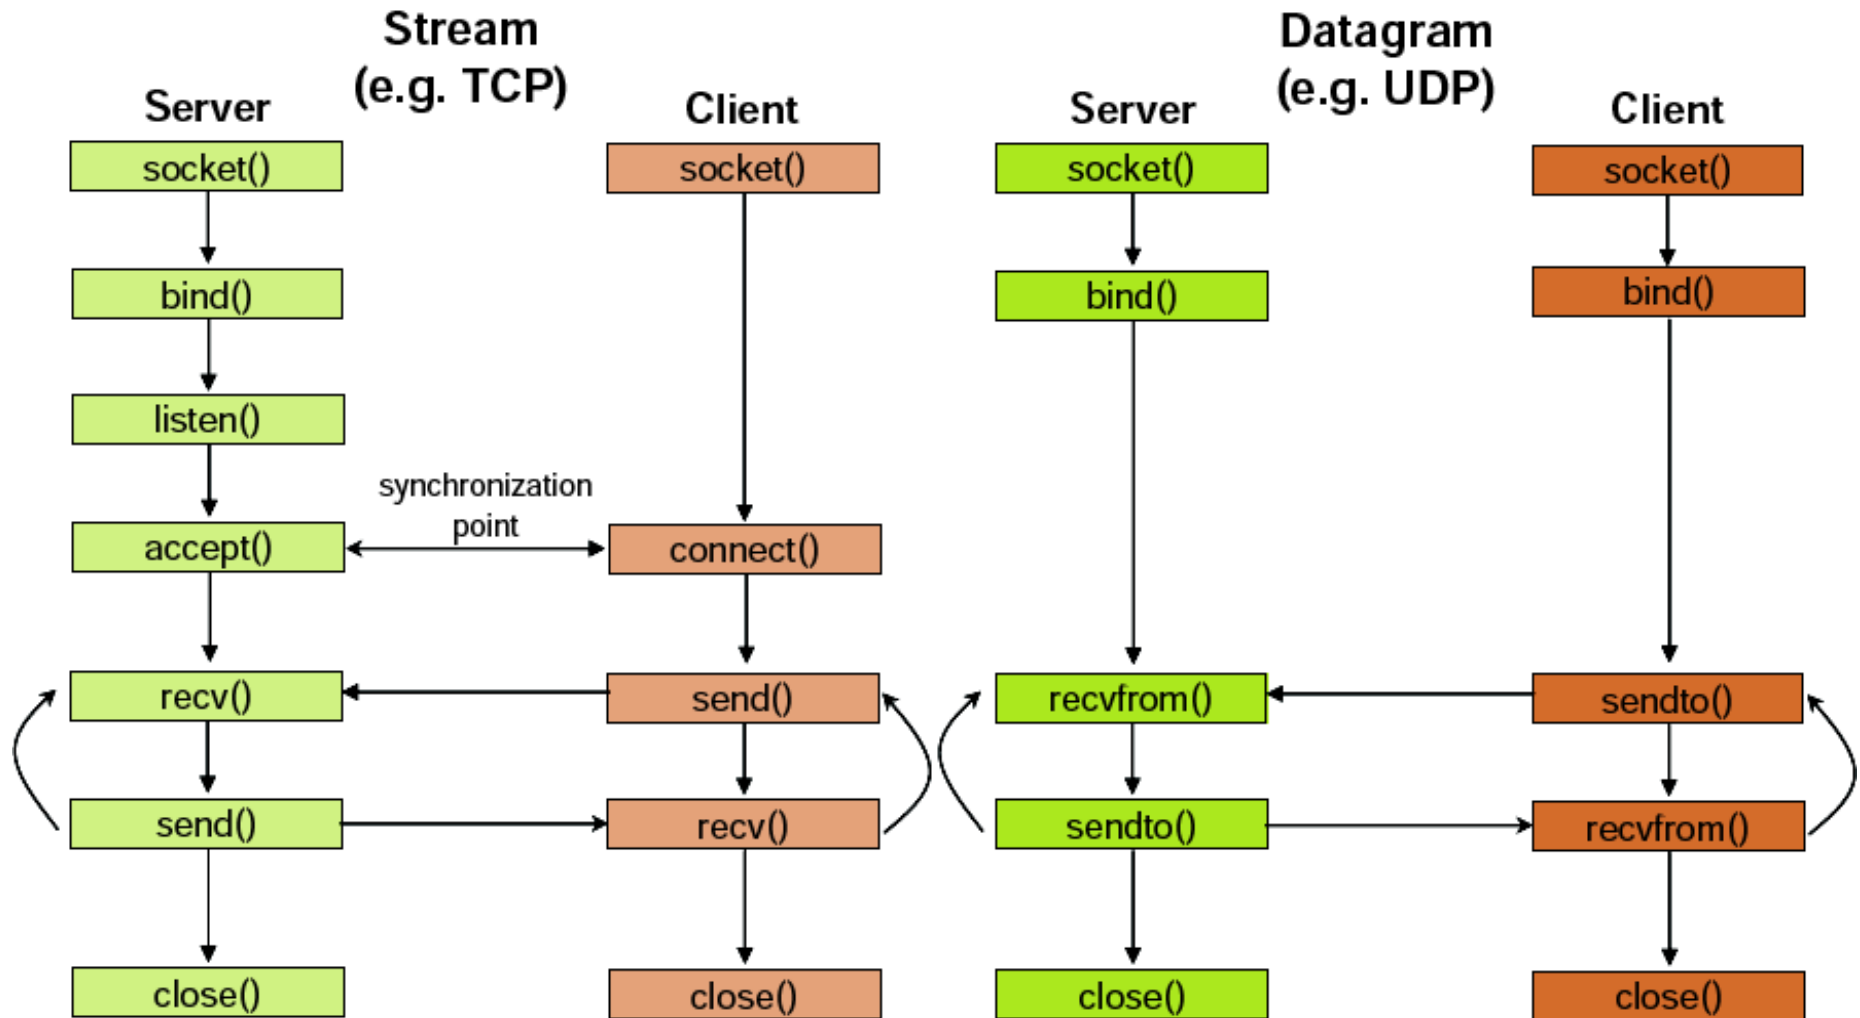
## Client

1. Create a TCP socket
2. Establish connection
3. Communicate
4. Close the connection

## Server

1. Create a TCP socket
2. Assign a port to socket
3. Set socket to listen
4. Repeatedly:
  - a. **Accept new connection**
  - b. Communicate
  - c. Close the connection

# Client - Server Communication - Unix



# Socket Options

- `getsockopt` and `setsockopt` allow socket options values to be queried and set, respectively
- `int getsockopt (sockid, level, optName, optVal, optLen);`
  - `sockid`: integer, socket descriptor
  - `level`: integer, the layers of the protocol stack (socket, TCP, IP)
  - `optName`: integer, option
  - `optVal`: pointer to a buffer; upon return it contains the value of the specified option
  - `optLen`: integer, in-out parameterit returns -1 if an error occurred
- `int setsockopt (sockid, level, optName, optVal, optLen);`
  - `optLen` is now only an input parameter



# Socket Options

## - Table

| <i>optName</i>           | Type          | Values        | Description                                                                                                           |
|--------------------------|---------------|---------------|-----------------------------------------------------------------------------------------------------------------------|
| <b>SOL_SOCKET Level</b>  |               |               |                                                                                                                       |
| SO_BROADCAST             | int           | 0,1           | Broadcast allowed                                                                                                     |
| SO_KEEPALIVE             | int           | 0,1           | Keepalive messages enabled (if implemented by the protocol)                                                           |
| SO_LINGER                | linger{}      | time          | Time to delay close() return waiting for confirmation (see Section 6.4.2)                                             |
| SO_RCVBUF                | int           | bytes         | Bytes in the socket receive buffer (see code on page 44 and Section 6.1)                                              |
| SO_RCVLOWAT              | int           | bytes         | Minimum number of available bytes that will cause recv() to return                                                    |
| SO_REUSEADDR             | int           | 0,1           | Binding allowed (under certain conditions) to an address or port already in use (see Section 6.4 and 6.5)             |
| SO_SNDLOWAT              | int           | bytes         | Minimum bytes to send a packet                                                                                        |
| SO_SNDBUF                | int           | bytes         | Bytes in the socket send buffer (see Section 6.1)                                                                     |
| <b>IPPROTO_TCP Level</b> |               |               |                                                                                                                       |
| TCP_MAX                  | int           | seconds       | Seconds between keepalive messages.                                                                                   |
| TCP_NODELAY              | int           | 0,1           | Disallow delay for data merging (Nagle's algorithm)                                                                   |
| <b>IPPROTO_IP Level</b>  |               |               |                                                                                                                       |
| IP_TTL                   | int           | 0-255         | Time-to-live for unicast IP packets                                                                                   |
| IP_MULTICAST_TTL         | unsigned char | 0-255         | Time-to-live for multicast IP packets (see MulticastSender.c on page 81)                                              |
| IP_MULTICAST_LOOP        | int           | 0,1           | Enables multicast socket to receive packets it sent                                                                   |
| IP_ADD_MEMBERSHIP        | ip_mreq{}     | group address | Enables reception of packets addressed to the specified multicast group (see MulticastReceiver.c on page 83)—set only |
| IP_DROP_MEMBERSHIP       | ip_mreq{}     | group address | Disables reception of packets addressed to the specified multicast group—set only                                     |

# Socket Options - Example

- Fetch and then double the current number of bytes in the socket's receive buffer

```
int rcvBufferSize;
int sockOptSize;
...
/* Retrieve and print the default buffer size */
sockOptSize = sizeof(rcvBuffSize);
if (getsockopt(sock, SOL_SOCKET, SO_RCVBUF, &rcvBufferSize, &sockOptSize) < 0)
 DieWithError("getsockopt() failed");
printf("Initial Receive Buffer Size: %d\n", rcvBufferSize);

/* Double the buffer size */
rcvBufferSize *= 2;

/* Set the buffer size to new value */
if (setsockopt(sock, SOL_SOCKET, SO_RCVBUF, &rcvBufferSize,
 sizeof(rcvBufferSize)) < 0)
 DieWithError("setsockopt() failed");
```



---

# Iterative Stream Socket Server

- Handles one client at a time
- Additional clients can connect while one is being served
  - connections are established
  - they are able to send requestsbut, the server will respond after it finishes with the first client
- 👍 Works well if each client required a small, bounded amount of work by the server
- 👎 otherwise, the clients experience long delays

# Iterative Server - Example: echo using stream socket

```
#include <stdio.h> /* for printf() and fprintf() */
#include <sys/socket.h> /* for socket(), bind(), connect(), recv() and send() */
#include <arpa/inet.h> /* for sockaddr_in and inet_ntoa() */
#include <stdlib.h> /* for atoi() and exit() */
#include <string.h> /* for memset() */
#include <unistd.h> /* for close() */

#define MAXPENDING 5 /* Maximum outstanding connection requests */

void DieWithError(char *errorMessage); /* Error handling function */
void HandleTCPClient(int clntSocket); /* TCP client handling function */

int main(int argc, char *argv[]) {
 int servSock; /* Socket descriptor for server */
 int clntSock; /* Socket descriptor for client */
 struct sockaddr_in echoServAddr; /* Local address */
 struct sockaddr_in echoClntAddr; /* Client address */
 unsigned short echoServPort; /* Server port */
 unsigned int clntLen; /* Length of client address data structure */

 if (argc != 2) { /* Test for correct number of arguments */
 fprintf(stderr, "Usage: %s <Server Port>\n", argv[0]);
 exit(1);
 }

 echoServPort = atoi(argv[1]); /* First arg: local port */

 /* Create socket for incoming connections */
 if ((servSock = socket(PF_INET, SOCK_STREAM, IPPROTO_TCP)) < 0)
 DieWithError("socket() failed");

 ...
}
```

# Iterative Server - Example: echo using stream socket

```
...

/* Construct local address structure */
memset(&echoServAddr, 0, sizeof(echoServAddr)); /* Zero out structure */
echoServAddr.sin_family = AF_INET; /* Internet address family */
echoServAddr.sin_addr.s_addr = htonl(INADDR_ANY); /* Any incoming interface */
echoServAddr.sin_port = htons(echoServPort); /* Local port */

/* Bind to the local address */
if (bind(servSock, (struct sockaddr *) &echoServAddr, sizeof(echoServAddr)) < 0)
 DieWithError("bind() failed");

/* Mark the socket so it will listen for incoming connections */
if (listen(servSock, MAXPENDING) < 0)
 DieWithError("listen() failed");

for (;;) /* Run forever */
{
 /* Set the size of the in-out parameter */
 clntLen = sizeof(echoClntAddr);

 /* Wait for a client to connect */
 if ((clntSock = accept(servSock, (struct sockaddr *) &echoClntAddr,
 &clntLen)) < 0)
 DieWithError("accept() failed");

 /* clntSock is connected to a client! */

 printf("Handling client %s\n", inet_ntoa(echoClntAddr.sin_addr));

 HandleTCPClient(clntSock);
}
/* NOT REACHED */
}
```

# Iterative Server - Example: echo using stream socket

```
#define RCVBUFSIZE 32 /* Size of receive buffer */

void HandleTCPClient(int clntSocket)
{
 char echoBuffer[RCVBUFSIZE]; /* Buffer for echo string */
 int recvMsgSize; /* Size of received message */

 /* Receive message from client */
 if ((recvMsgSize = recv(clntSocket, echoBuffer, RCVBUFSIZE, 0)) < 0)
 DieWithError("recv() failed");

 /* Send received string and receive again until end of transmission */
 while (recvMsgSize > 0) /* zero indicates end of transmission */
 {
 /* Echo message back to client */
 if (send(clntSocket, echoBuffer, recvMsgSize, 0) != recvMsgSize)
 DieWithError("send() failed");

 /* See if there is more data to receive */
 if ((recvMsgSize = recv(clntSocket, echoBuffer, RCVBUFSIZE, 0)) < 0)
 DieWithError("recv() failed");
 }

 close(clntSocket); /* Close client socket */
}
```

---

# Multitasking - Per-Client Process

- For each client connection request, a new process is created to handle the communication
- `int fork();`
  - a new process is created, identical to the calling process, except for its process ID and the return value it receives from `fork()`
  - returns 0 to **child** process, and the process ID of the new child to **parent**

## Caution:

- when a child process terminates, it does not automatically disappear
- use `waitpid()` to parent in order to “harvest” zombies

# Multitasking - Per-Client Process

- Example: echo using stream socket

```
#include <sys/wait.h> /* for waitpid() */

int main(int argc, char *argv[]) {
 int servSock; /* Socket descriptor for server */
 int clntSock; /* Socket descriptor for client */
 unsigned short echoServPort; /* Server port */
 pid_t processID; /* Process ID from fork() */
 unsigned int childProcCount = 0; /* Number of child processes */

 if (argc != 2) { /* Test for correct number of arguments */
 fprintf(stderr, "Usage: %s <Server Port>\n", argv[0]);
 exit(1);
 }
 echoServPort = atoi(argv[1]); /* First arg: local port */

 servSock = CreateTCPServerSocket(echoServPort);

 for (;;) { /* Run forever */
 clntSock = AcceptTCPConnection(servSock);

 if ((processID = fork()) < 0) DieWithError ("fork() failed"); /* Fork child process */
 else if (processID == 0) { /* This is the child process */
 close(servSock); /* child closes listening socket */
 HandleTCPClient(clntSock);
 exit(0); /* child process terminates */
 }

 close(clntSock); /* parent closes child socket */
 childProcCount++; /* Increment number of outstanding child processes */

 ...
 }
}
```

# Multitasking - Per-Client Process

- Example: echo using stream socket

```
...
while (childProcCount) {
 processID = waitpid((pid_t) -1, NULL, WHOANG);
 if (processID < 0) DieWithError ("...");
 else if (processID == 0) break;
 else childProcCount--;
}
/* NOT REACHED */
}
```

---

# Multitasking - Per-Client Thread

- 👎 Forking a new process is expensive
  - ❑ duplicate the entire state (memory, stack, file/socket descriptors, ...)
- 👍 Threads decrease this cost by allowing multitasking within the same process
  - ❑ threads share the same address space (code and data)

An example is provided using POSIX Threads



# Multitasking - Per-Client Thread

- Example: echo using stream socket

```
#include <pthread.h> /* for POSIX threads */

void *ThreadMain(void *arg) /* Main program of a thread */

struct ThreadArgs { /* Structure of arguments to pass to client thread */
 int clntSock; /* socket descriptor for client */
};

int main(int argc, char *argv[]) {
 int servSock; /* Socket descriptor for server */
 int clntSock; /* Socket descriptor for client */
 unsigned short echoServPort; /* Server port */
 pthread_t threadID; /* Thread ID from pthread_create() */
 struct ThreadArgs *threadArgs; /* Pointer to argument structure for thread */

 if (argc != 2) { /* Test for correct number of arguments */
 fprintf(stderr, "Usage: %s <Server Port>\n", argv[0]);
 exit(1);
 }
 echoServPort = atoi(argv[1]); /* First arg: local port */

 servSock = CreateTCPServerSocket(echoServPort);

 for (;;) { /* Run forever */
 clntSock = AcceptTCPConnection(servSock);

 /* Create separate memory for client argument */
 if ((threadArgs = (struct ThreadArgs *) malloc(sizeof(struct ThreadArgs))) == NULL) DieWithError("...");
 threadArgs->clntSock = clntSock;

 /* Create client thread */
 if (pthread_create (&threadID, NULL, ThreadMain, (void *) threadArgs) != 0) DieWithError("...");
 }
 /* NOT REACHED */
}
```

# Multitasking - Per-Client Thread

- Example: echo using stream socket

```
void *ThreadMain(void *threadArgs)
{
 int clntSock; /* Socket descriptor for client connection */

 pthread_detach(pthread_self()); /* Guarantees that thread resources are deallocated upon return */

 /* Extract socket file descriptor from argument */
 clntSock = ((struct ThreadArgs *) threadArgs) -> clntSock;
 free(threadArgs); /* Deallocate memory for argument */

 HandleTCPClient(clntSock);

 return (NULL);
}
```

---

# Multitasking - Constrained

- Both process and thread incurs **overhead**
  - creation, scheduling and context switching
- As their numbers increases
  - this overhead increases
  - after some point it would be better if a client was blocked
- Solution: **Constrained multitasking**. The server:
  - begins, creating, binding and listening to a socket
  - creates a number of processes, each loops forever and accept connections from the same socket
  - when a connection is established
    - the client socket descriptor is returned to only one process
    - the other remain blocked

# Multitasking - Constrained

## - Example: echo using stream socket

```
void ProcessMain(int servSock); /* Main program of process */

int main(int argc, char *argv[]) {
 int servSock; /* Socket descriptor for server*/
 unsigned short echoServPort; /* Server port */
 pid_t processID; /* Process ID */
 unsigned int processLimit; /* Number of child processes to create */
 unsigned int processCt; /* Process counter */

 if (argc != 3) { /* Test for correct number of arguments */
 fprintf(stderr, "Usage: %s <SERVER PORT> <FORK LIMIT>\n", argv[0]);
 exit(1);
 }

 echoServPort = atoi(argv[1]); /* First arg: local port */
 processLimit = atoi(argv[2]); /* Second arg: number of child processes */

 servSock = CreateTCPServerSocket(echoServPort);

 for (processCt=0; processCt < processLimit; processCt++)
 if ((processID = fork()) < 0) DieWithError("fork() failed"); /* Fork child process */
 else if (processID == 0) ProcessMain(servSock); /* If this is the child process */

 exit(0); /* The children will carry on */
}

void ProcessMain(int servSock) {
 int clntSock; /* Socket descriptor for client connection */

 for (;;) { /* Run forever */
 clntSock = AcceptTCPConnection(servSock);
 printf("with child process: %d\n", (unsigned int) getpid());
 HandleTCPClient(clntSock);
 }
}
```

---

# The End - Questions

